


Multiple Layer Perception

A multi-layer perceptron is class of ANN that consists of multiple layer (hidden layers) of neurons in a feed-forward network.

Architecture of MLP

at least (min^m) one hidden layer is present in MLP.






- receives the input data (after flattening) and connect each of the input features (columns) to neurons in the input layer

- one or more hidden layer(s) between input and output layers where the actual learning happens-

- produces the final prediction

- Regression
- Classification
 - Binary
 - Multi-class.




- No learning or calculation

Neural Network Terminologies

Fully Connected Network

- a layer where each and every neuron is connected to every neuron in the preceding previous layer.

Neural Network Terminologies




Single Layer Perceptron (SLP) vs Multi Layer Perceptron (MLP)		
Feature	SLP (Single Layer Perceptron)	MLP (Multi Layer Perceptron)
Layers	1 layer (input → output)	2 or more layers (input → hidden(s) → output)
Neurons	No hidden layer, just output neuron(s)	One or more hidden layers with multiple neurons
Functions	Can only solve linearly separable problems	Can solve non-linear and more complex problems
Learning	Simple weights & bias update (perceptron rule)	Uses backpropagation and more complex optimization
Representation	Linear decision boundaries	Can learn complex, non-linear boundaries
Use cases	Very basic classification tasks	Most modern neural network applications

Computation: simple and fast
(but just a proof of concept) more computationally intensive than SLP
(state of art NN model)


Working of MLP


Step #1 Getting input data ready.

- to classify a handwritten digit



Step #2 Take (weighted sum of inputs + bias) \rightarrow in hidden layer





$HL\#1$



Each neuron in the 1^{st} HL receives a weighted sum of inputs or features (x_1, x_2, x_3)
 \oplus
 from the input layer.
 (bias)

Computing the weighted sum of inputs along with bias:

$$Z_j^l = W_{ij}^l x_i + b_j^l$$


where



Z_j^l : is the weight sum for neurons j in the layer l

W_{ij}^l : is the weight between neuron i from the previous layer ($l-1$)
 and neuron j in the current layer (l)


x_i : is the input from neuron i in the previous layer


b_j^l : is the bias associated with neuron j in the layer l





$$\begin{aligned}
 z_1^{(l)} &= w_{11}^{(l)} x_1 + w_{21}^{(l)} x_2 + w_{31}^{(l)} x_3 + b_1^{(l)} \leftarrow 1^{\text{st}} \text{neuron } | l=1 \\
 z_2^{(l)} &= w_{12}^{(l)} x_1 + w_{22}^{(l)} x_2 + w_{32}^{(l)} x_3 + b_2^{(l)} \leftarrow 2^{\text{nd}} \text{neuron } | l=1 \\
 z_3^{(l)} &= w_{13}^{(l)} x_1 + w_{23}^{(l)} x_2 + w_{33}^{(l)} x_3 + b_3^{(l)} \leftarrow 3^{\text{rd}} \text{neuron } | l=1 \\
 z_4^{(l)} &= w_{14}^{(l)} x_1 + w_{24}^{(l)} x_2 + w_{34}^{(l)} x_3 + b_4^{(l)} \leftarrow 4^{\text{th}} \text{neuron } | l=1
 \end{aligned}$$





Input vector: $X = \begin{bmatrix} x_1 \\ x_2 \\ x_3 \end{bmatrix}_{3 \times 1}$

Bias vector: $b = \begin{bmatrix} b_1 \\ b_2 \\ b_3 \\ b_4 \end{bmatrix}_{4 \times 1}$

layer #2.

$$Z_l = W * X + b = [Wx] + b$$

$4 \times 3 * 3 \times 1$
 4×1

4×1


addition is valid.

$$Z^{(l)} = \begin{bmatrix} w_{11} & w_{21} & w_{31} \\ w_{12} & w_{22} & w_{32} \\ w_{13} & w_{23} & w_{33} \\ w_{14} & w_{24} & w_{34} \end{bmatrix} * \begin{bmatrix} x_1 \\ x_2 \\ x_3 \end{bmatrix} + \begin{bmatrix} b_1 \\ b_2 \\ b_3 \\ b_4 \end{bmatrix}$$

for layer = l.

$$z_1 = w_{11}x_1 + w_{21}x_2 + w_{31}x_3 + b_1$$

$$z_3 = w_{13}x_1 + w_{23}x_2 + w_{33}x_3 + b_3$$





Task: Update the equations for $l=2$ and $l=3$.


- Manish Kaushik ✓

For layer = l and 1st neuron ($j=1$)

$$z_1^{(l)} = w_{11}^{(l)}x_1 + w_{21}^{(l)}x_2 + w_{31}^{(l)}x_3 + b_1^{(l)} \leftarrow \text{at neuron } l=1$$



Activation Functions



Purpose: Activation function introduces **non-linearity** into the neural net so it can learn complex patterns (otherwise, the net would behave like a linear regression model regardless of its depth (many hidden layers))

1. Sigmoid Activation Function


$$f(z) = \sigma(z) = \frac{1}{1+e^{-z}}$$

: primarily used for binary classification model.


$$z=0; \sigma(0) = \frac{1}{1+e^0} = \frac{1}{1+1} = \frac{1}{2} = 0.5$$

$$z \rightarrow \infty; \sigma(\infty) = \frac{1}{1+e^{-\infty}} = 1$$

$$z \rightarrow -\infty; \sigma(-\infty) = \frac{1}{1+e^{(\infty)}} = \frac{1}{1+e^\infty} = \frac{1}{\infty} \rightarrow 0$$



$$\begin{cases} z \geq 0.5 \rightarrow \text{class } 1 \\ z < 0.5 \rightarrow \text{class } 0 \end{cases}$$



Note: Technically speaking, a sigmoid is **any S-shaped curve** that flattens out near its minimum and maximum values.

tanh: hyperbolic tangent
 ↳ RNN / LSTM → will be discussed here.

Range of sigmoid activation function: $(0, 1)$

Application: Binary classification \rightarrow (output layer) \rightarrow probabilistic outputs

Pros:

- smooth gradient \rightarrow sigmoid is differentiable function
- and range is very useful for probabilities $(0, 1)$


Cons: $\text{pro-tip} \rightarrow (\text{RNN}[\text{LSTM}])$

- Vanishing Gradient: \rightarrow is a problem for very large/small input values
(will be discussed in RNN)
- Not zero-centered (all outputs are going to be positive)
 \hookrightarrow which eventually leads to slow convergence.

FYI


2. ReLU: Rectified Linear Unit


Linear Activation Function




- output is proportional to the input
- lines pass through the origin $\rightarrow y = mx$ form
- since, lines are passing through origin its usage is limited in deep learning given that it lacks non-linearity

Rectified Linear Unit





* pro-tip Why is it called **rectified**??



Electronics Coach

It's called 'rectified' because of how it fixes or clips the -ve values to zero, just like a rectifier in electronics engg. → which only allows positive parts of a signal to pass.

Similarly, ReLU activation functions rectifies the input by

- Keeping all positive values
- and replacing all negative values with zero.

$$\text{ReLU}(z) = \max(0, z)$$


In some sense, ReLU is used as activation function to introduce non-linearity but at the same time it also adds **sparsity** *

What is sparse data?

Sparse data or sparsity is an important concept in both ML and DL specially when dealing with large datasets.

- Sparse data means most of the values in data are zero or entry.

NETFLIX



- In such datasets, most of the entries are going to be '0' → sparse data.

Note: Text, recommendations, some computer vision problems, NLP tasks naturally produce sparse data.

ReLU adds non-linearity by cutting off the -ve values and output becomes sparse.



many neurons output '0' values

Given a neuron's (weighted sum of inputs + bias) becomes -ve for all inputs → ReLU will map it to '0'.



Dying 'ReLU' problem



Analogy → A person getting brain stroke.

① abc → Dying ReLU!

Types of PARALYSIS



Two Types of Stroke



→ (Ray of hope → Physiotherapy)

Leaky ReLU

Dying ReLU Problem

- If a neuron's input ($Wx + b$) becomes negative, ReLU outputs 0
- During backpropagation, the gradient of ReLU is also 0 for negative inputs

$$f'(z) = \begin{cases} 1 & z > 0 \\ 0 & z \leq 0 \end{cases}$$

(Neuron)

→ it stops updating as no gradient and it stays dead and always output '0' forever.

Example:

If a neuron's weights along with bias produce

$$z = -3 \text{ then}$$

$\left\{ \begin{array}{l} \text{output: } f(z) = 0 \\ \text{gradient: } f'(z) = 0 \end{array} \right.$
 → No weight update during backpropagation

- The neuron is effectively dead — permanently inactive.

Pros of using ReLU

Even though neurons can be dead while using ReLU as AF and can lead to "dying ReLU problem" however it still the most popular activation functions in hidden layers because:

✓ # Simplicity

→ Mathematical formulation: $\max(0, z) \rightarrow$ very cheap to compute ↓


- # Simplicity → Mathematical formulation: $\max(0, z) \rightarrow$ very cheap to compute
computational cost is less
- # Faster Convergence → ReLU doesn't saturate for large positive values → gradients stay/keep changing leading to faster convergence.
- # Sparse activations → ReLU makes some of the neurons dead or '0' → leads to efficient computations of the neural network model

Cons of using ReLU

- # Dying ReLU → it is a problem which refers to the scenario where a significant number of neurons in the neural net model always outputs zero after applying the ReLU AF.
- # it is also non-zero centered just like sigmoid function.

Will do it later (module #3)

Vishwas: Although ReLU has drawbacks such as producing biased (non-zero-centered) outputs, these issues can be effectively taken care of by using good weight initialization (e.g., Kaiming) and Batch Normalization. Because of this ReLU is one of the most widely used activation functions for hidden layers in deep neural networks.



OLED: Organic Light Emitting Diode

