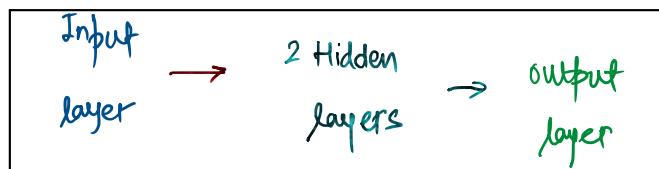
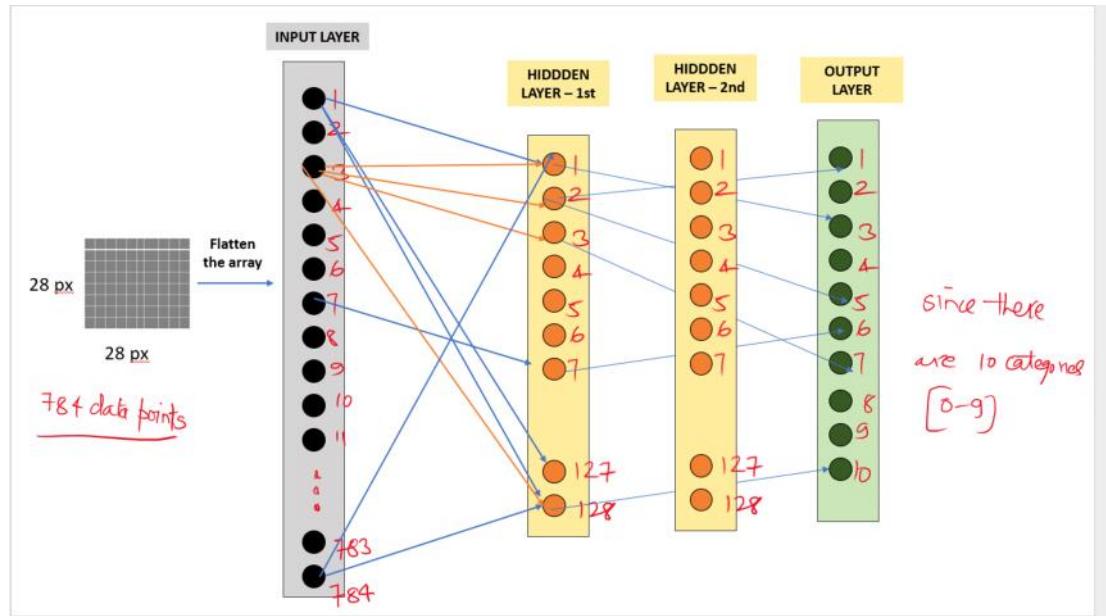


MNIST Code Explanation

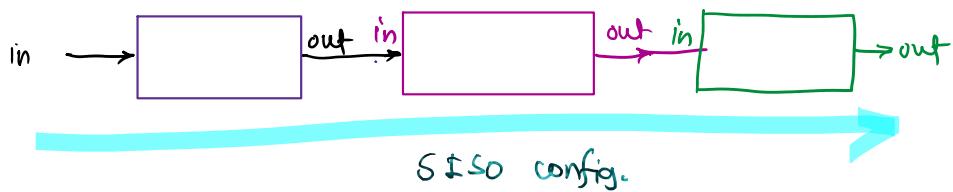
01 November 2025 12:24



```
: model = tf.keras.models.Sequential() # creates the model object of the class `Sequential` which is basically a linear stack of layers
: model
    ↗ creates an empty sequential model object in TF/Keras
    ↗ is a linear stack of layers
        ↓
        to add one layer at a time using .add()
```

Each layer has exactly **one input** and **one output** Tensor

SISO: single input single output



```
### Input Layer
```

Raw Data for MNIST

```

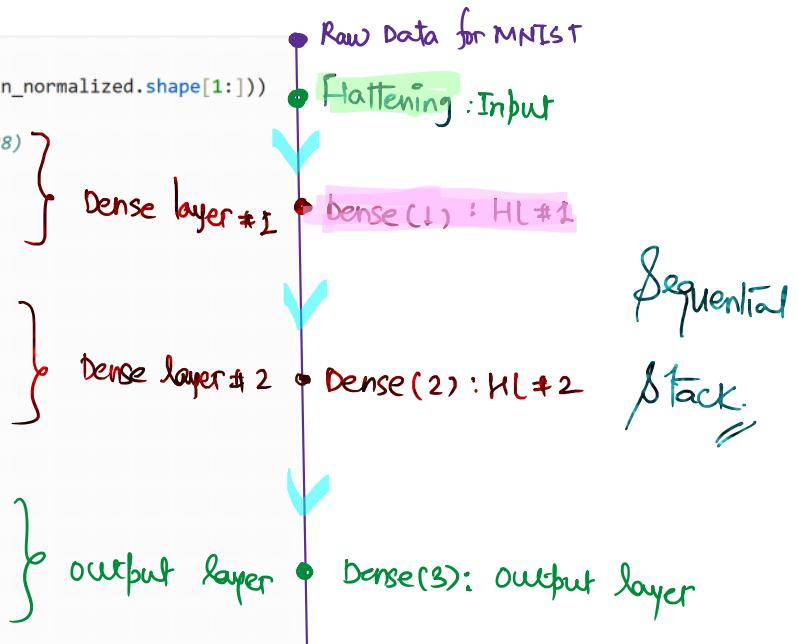
    : ### Input Layer
    model.add(tf.keras.layers.Flatten(input_shape = x_train_normalized.shape[1:]))

    ### First Hidden Layer having size (no. of neurons = 28)
    model.add(tf.keras.layers.Dense(
        units = 128,
        activation=tf.keras.activations.relu
    ))

    ### Second Hidden Layer having size 128
    model.add(tf.keras.layers.Dense(
        units = 128,
        activation=tf.keras.activations.relu
    ))

    ### Output Layer
    model.add(tf.keras.layers.Dense(
        units=10,
        activation = tf.keras.activations.softmax
    ))

```



First Hidden layer

```

    ### First Hidden Layer having size (no. of neurons = 28)
    model.add(tf.keras.layers.Dense(
        units = 128,
        activation=tf.keras.activations.relu
    ))

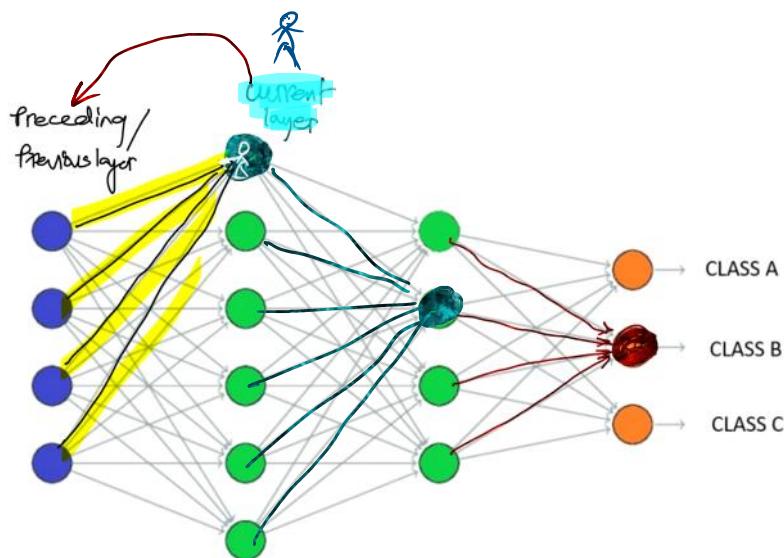
```

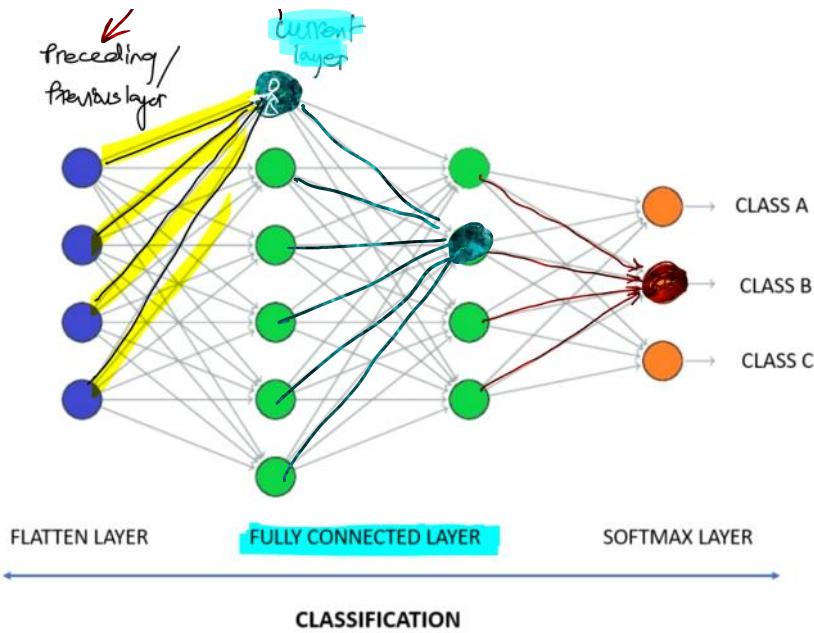
} Dense layer #1

Pro tip What is dense layer in a neural network model?

A dense layer is also known as fully connected layer
is a fundamental building block in ANN.

In this layer, each neuron in the current layer is
connected to every neuron in the preceding / previous
layer





https://www.tensorflow.org/api_docs/python/tf/keras/layers/Dense

```
tf.keras.layers.Dense(
    units,           → no. of neurons
    activation=None, → RelU, softmax, sigmoid etc.
    use_bias=True,   → by default → TRUE
    kernel_initializer='glorot_uniform',
    bias_initializer='zeros',   → a technique to initialize weights
    kernel_regularizer=None,   → bias initialized with 0 values.
    bias_regularizer=None,
    activity_regularizer=None,
    kernel_constraint=None,
    bias_constraint=None,
    lora_rank=None,
    **kwargs
)
```

Annotations on the code:

- 'units' → no. of neurons
- 'activation=None' → RelU, softmax, sigmoid etc.
- 'use_bias=True' → by default → TRUE
- 'kernel_initializer='glorot_uniform'' → a technique to initialize weights
- 'bias_initializer='zeros'' → bias initialized with 0 values.
- 'kernel_regularizer=None' → adds a penalty to the layer weights during training (L_1/L_2) and helps prevent overfitting -
- 'bias_regularizer=None'
- 'activity_regularizer=None'
- 'kernel_constraint=None'
- 'bias_constraint=None'
- 'lora_rank=None'
- '**kwargs'

Task: Read about 'glorot-uniform'

Args	
units	Positive integer, dimensionality of the output space.
activation	Activation function to use. If you don't specify anything, no activation is applied (ie. "linear" activation: $a(x) = x$).
use_bias	Boolean, whether the layer uses a bias vector. X
kernel_initializer	Initializer for the kernel weights matrix.
bias_initializer	Initializer for the bias vector.
kernel_regularizer	Regularizer function applied to the kernel weights matrix.
bias_regularizer	Regularizer function applied to the bias vector.
activity_regularizer	Regularizer function applied to the output of the layer (its "activation").
kernel_constraint	Constraint function applied to the kernel weights matrix.
bias_constraint	Constraint function applied to the bias vector.
lora_rank	Optional integer. If set, the layer's forward pass will implement LoRA (Low-Rank Adaptation) with the provided rank. LoRA sets the layer's kernel to non-trainable and replaces it with a delta over the original kernel, obtained via multiplying two lower-rank trainable matrices. This can be useful to reduce the computation cost of fine-tuning large dense layers. You can also enable LoRA on an existing Dense layer by calling <code>layer.enable_lora(rank)</code> .

`model.summary()` → To generate a tabular structure of NN model's architecture along with Model: "sequential_3" output and # parameters.

Layer (type)	Output Shape	Param #
flatten_6 (Flatten)	(None, 784)	28 * 28 = 784 0
dense_9 (Dense)	(None, 128)	100,480
dense_10 (Dense)	(None, 128)	16,512
dense_11 (Dense)	(None, 10)	1,290

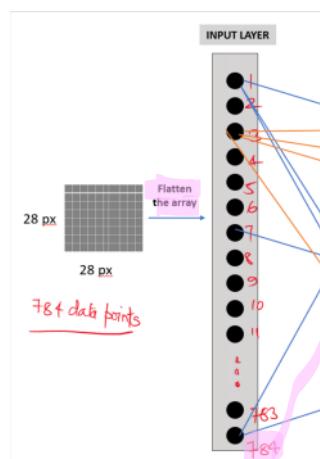
Total params: 118,282 (462.04 KB)

Trainable params: 118,282 (462.04 KB)

Non-trainable params: 0 (0.00 B)

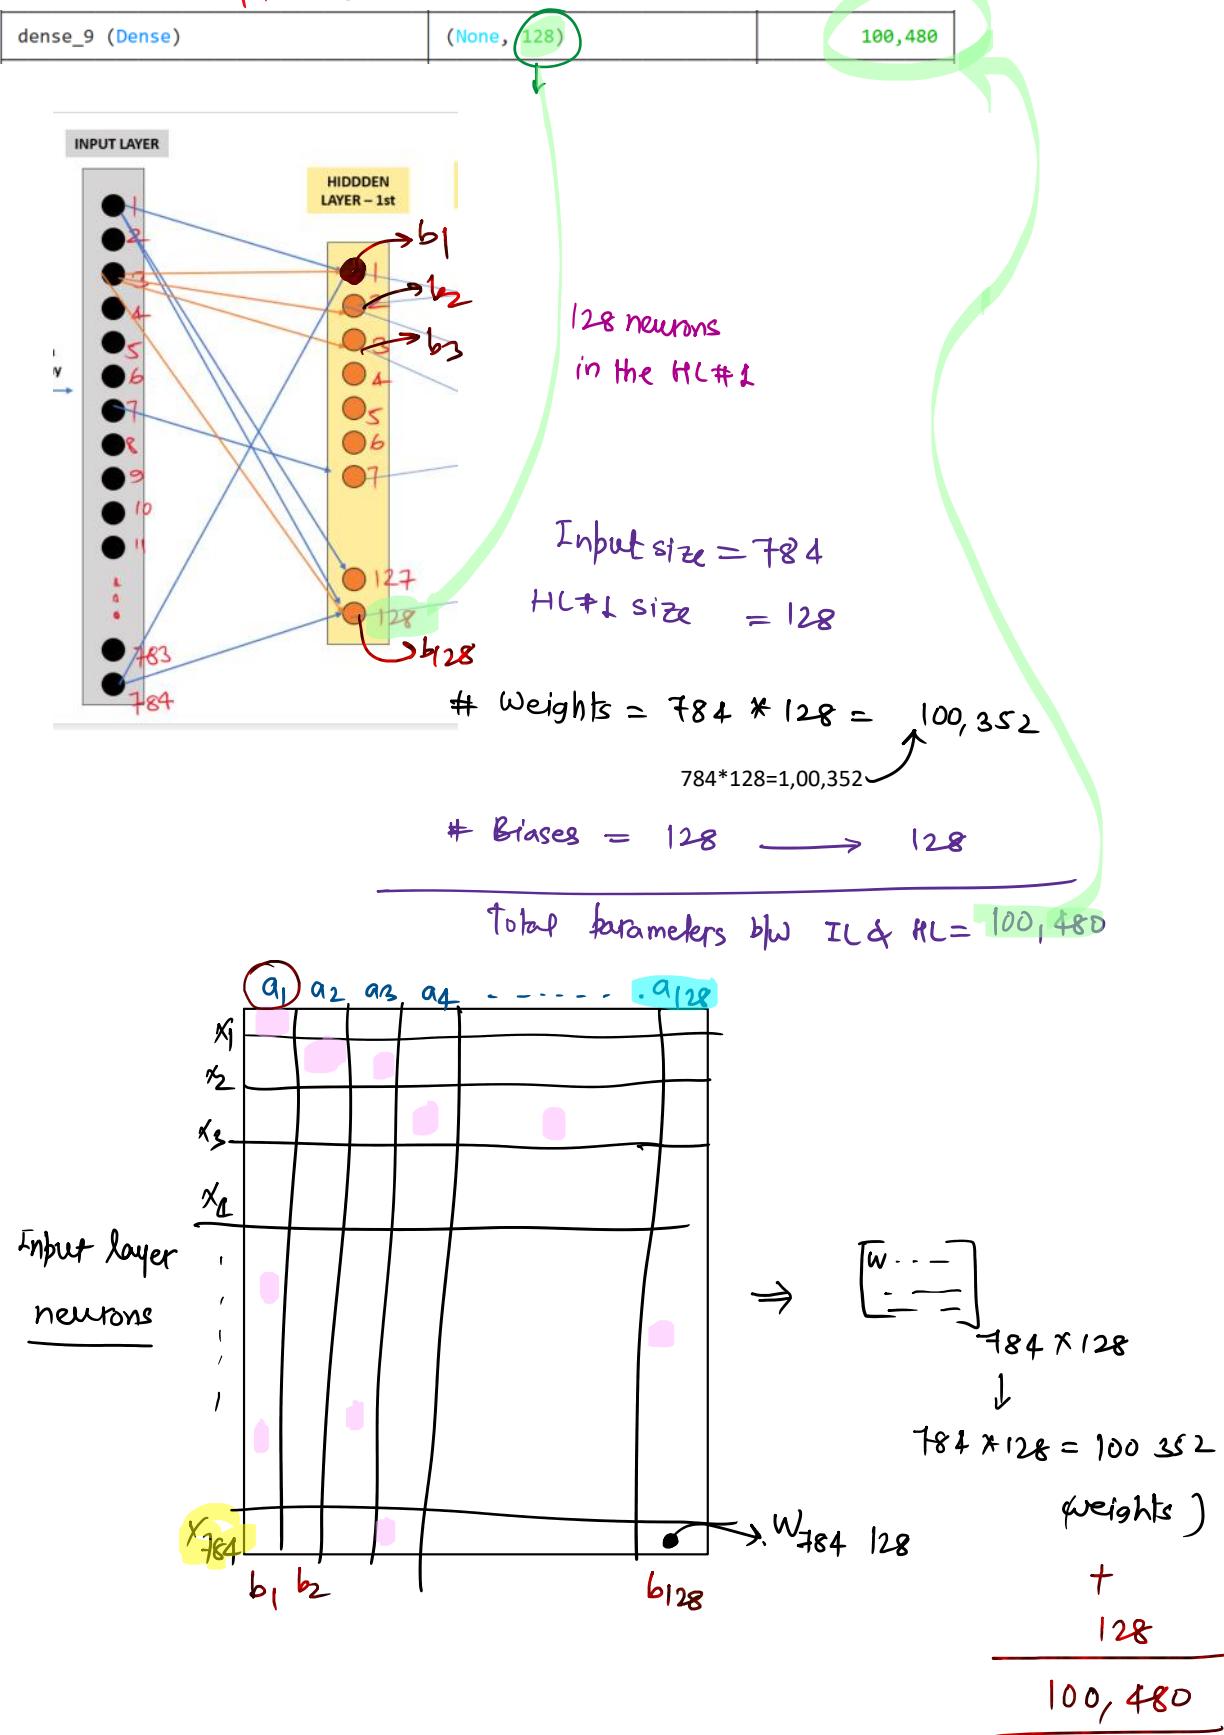
→ no learnable parameters
(No weights or bias)
In IL

batch-size can vary
↓
model doesn't know how many samples
it will get at once during training



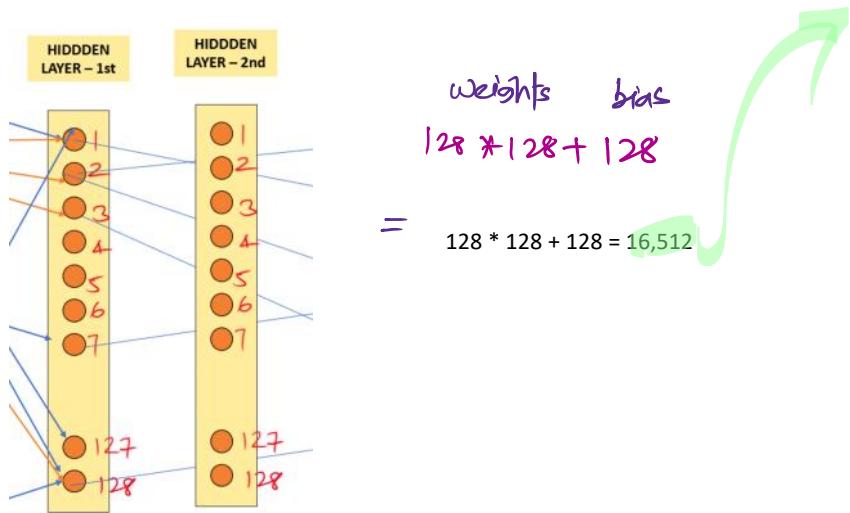
First Hidden Layer HL#1

First Hidden Layer HL#1



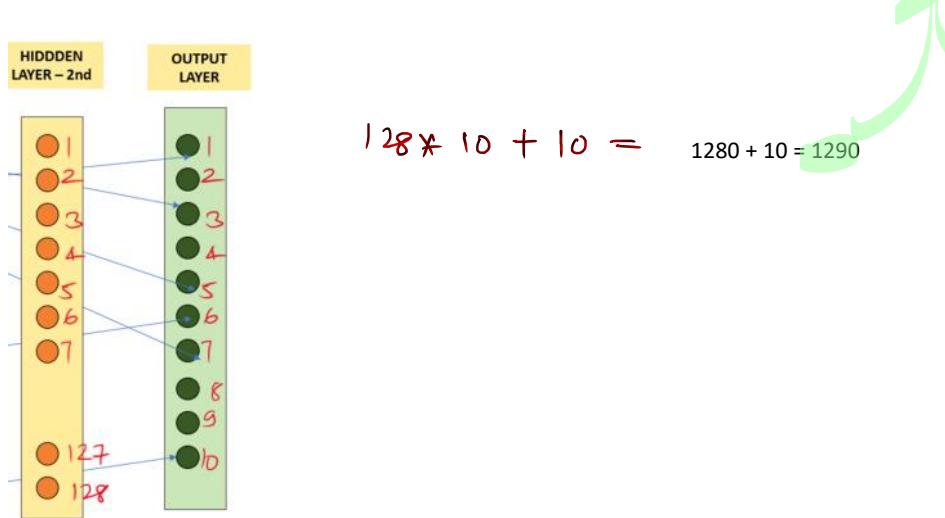
2nd Hidden layer





output layer

dense_11 (Dense)	(None, 10)	1,290
------------------	------------	-------



Total params: 118,282 (462.04 KB) ✓

Trainable params: 118,282 (462.04 KB) ✓

→ Non-trainable params: 0 (0.00 B)

118,282 parameters

1 byte = 8 bits

dtype: float32 by default

1 param → 32 bits → 4 bytes

approximate memory size required to store all model parameters (weights + biases)

d type : float32 by default

1 param \rightarrow 32 bits \rightarrow 4 bytes

$118282 \times 4 \text{ bytes}$ $118282 \times 4 = 473128 \text{ bytes}$

$$\frac{473128}{1024} = 462.0391 \text{ KB}$$

Note: 462.04 KB is NOT the full model size on the disk.

rather it's the size of model's weights only ~~size~~

Compile the model

```
[1]: adam_optimizer = tf.keras.optimizers.Adam(learning_rate=0.01)
[2]: model.compile(
    optimizer = adam_optimizer,
    loss = tf.keras.losses.sparse_categorical_crossentropy,
    metrics = ['accuracy'])
```

will discuss later

a popular choice for an optimizer

`model.compile()` \rightarrow model is being configured for the training by invoking

- appropriate optimizer
- loss function for multi-class
- metrics set to 'accuracy'

↳ given that MNIST is a balanced dataset.

both training and validation simultaneously for epoch.

```
### Train and validate the model [x] 784*60000
training_hist = model.fit(x_train_normalized, y_train, epochs = 30, validation_data=(x_test_normalized, y_test))

Signature:
model.fit(
    x=None,
    y=None,
    batch_size=None,
```

TF keras does both training and validation once per epoch.

```

model.fit(
    x=None,
    y=None,
    batch_size=None,
    epochs=1, → by default → epoch # 1
    verbose='auto',
    callbacks=None,
    validation_split=0.0,
    validation_data=None,
    shuffle=True,
    class_weight=None,
    sample_weight=None,
    initial_epoch=0,
    steps_per_epoch=None,
    validation_steps=None,
    validation_batch_size=None,
    validation_freq=1,
)

```

batch_size: Integer or 'None'.
Number of samples per gradient update.
If unspecified, 'batch_size' will default to 32.
Do not specify the 'batch_size' if your input data 'x' is a 'keras.utils.PyDataset', 'tf.data.Dataset', 'torch.utils.data.DataLoader' or Python generator function since they generate batches.

output/epoch total no. of epochs # 30

Epoch 1/30
1875/1875 → 6s 2ms/step - accuracy: 0.8821 - loss: 0.3846 - val_accuracy: 0.9488 - val_loss: 0.1873

$$\text{mini-batches} = 60000/32 = 1875$$

1/1875

2/1875

:

1875/1875

Epoch # (1/30) → indicates the training process by epoch
- model is currently in epoch # 1 out of total 30 epochs.

1875/1875 : Indicates no. of training batches/epoch.

6s 2ms/step : it indicates epoch # 1 took ~ 6s seconds in total and 2ms on an avg. per batch.

Training

- accuracy: 0.8821 - loss: 0.3846

Per epoch # 1

$$\text{Training accuracy} = 88.21\%$$

$$\text{Training loss} = 0.3846$$

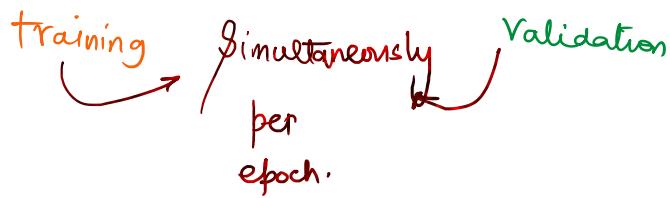
Testing

- val_accuracy: 0.9488 - val_loss: 0.1873

$$\text{Validation accuracy} = 94.88\%$$

$$\text{validation loss} = 0.1873$$

```
### Train and validate the model  
training_hist = model.fit(x_train_normalized, y_train, epochs = 30, validation_data=(x_test_normalized, y_test))
```



Note: Default behavior of `model.fit()`

However you can run all the epochs only for training and then run validation only once.

option #1 Train without validation during epochs.

`model.fit(x=x_train_normalized, y=y_train, epochs=30, verbose=1)`

Run validation block only once at the end

`val_loss, val_acc = model.evaluate(x=x_test_normalized, y=y_test, verbose=1)`